

# Python scripting for network analysis — III

Lab session on March 5th

# Homework review

# Homework: ring & shortcuts network

We construct the following network:

1. ring of  $N$  nodes with first and second neighbor connections
2. add  $N/2$  shortcuts (additional edges) at random

Question: what is the average clustering coefficient of this network as  $N \rightarrow \infty$ ?

1. Generate network
2. Measure quantity (average clustering)
3. Estimate for  $N \rightarrow \infty$

# Generate network

1. Start with ring
2. Add shortcuts

```
In [2]: import networkx
import matplotlib.pyplot as plt
```

```
def ring_network(N):
    g = networkx.Graph()
    for i in range(N):
        g.add_edge(i, (i+1)%N)
        g.add_edge(i, (i+2)%N)
    return g
```

```
In [3]: import random
def add_shortcuts_first_try(graph):
    nodes = list(graph.nodes())
    for i in range(int(len(nodes)/2)):
        node_a = random.choice(nodes)
        node_b = random.choice(nodes)
        graph.add_edge(node_a, node_b)
    return graph
```

```
In [4]: for i in range(10):
        g = add_shortcuts_first_try(ring_network(10))
        print(g.number_of_nodes(), g.number_of_edges())
```

```
10 24
10 22
10 21
10 23
10 22
10 23
10 24
10 22
10 24
10 21
```

```
In [5]: import random
def add_shortcuts(graph):
    if graph.number_of_nodes() < 6:
        raise ValueError("network is too small!")
    nodes = list(graph.nodes())
    for i in range(int(len(nodes)/2)):
        node_a = random.choice(nodes)
        node_b = random.choice(nodes)
        while graph.has_edge(node_a, node_b) or node_a == node_b:
            node_a = random.choice(nodes)
            node_b = random.choice(nodes)
        graph.add_edge(node_a, node_b)
    return graph

for i in range(10):
    g = add_shortcuts(ring_network(10))
    print(g.number_of_nodes(), g.number_of_edges())
```

```
10 25
10 25
10 25
10 25
10 25
10 25
10 25
10 25
10 25
10 25
```

# Estimating $N \rightarrow \infty$

```
In [6]: networkx.average_clustering(add_shortcuts(ring_network(10)))
```

```
Out[6]: 0.5809523809523809
```

```
In [7]: networkx.average_clustering(add_shortcuts(ring_network(100)))
```

```
Out[7]: 0.35037301587301584
```

```
In [8]: networkx.average_clustering(add_shortcuts(ring_network(100000)))
```


```
Out[8]: 0.34191382900428413
```



# Make convincing!

```
In [9]: N_values = [10,20,50,100,200,500, 1000, 2000,5000]
        values = [networkx.average_clustering(add_shortcuts(ring_network(N))) for N in N_v
        alues]
        plt.plot(N_values, values)
```

```
Out[9]: [<matplotlib.lines.Line2D at 0x7fd911baeb00>]
```



In [10]: values

Out[10]: [0.5733333333333335,  
0.446904761904762,  
0.36000000000000002,  
0.3610714285714285,  
0.3417619047619043,  
0.3446460317460319,  
0.340650000000000134,  
0.3421999999999986,  
0.34251857142856473]

# State the result!

- in this case, for example: "the average clustering as  $N \rightarrow \infty$  is between 0.340 and 0.342"
- specifying interval allows showing the uncertainty of the estimate

# Does our result match the theoretical estimate?

- theoretical estimate was 0.3
- this is outside the range of our estimate ("between 0.340 and 0.342"), thus the two results don't match

# What causes the discrepancy?

- Is it the approximation done for the theoretical estimate? (that we ignore triangles created by the shortcut edges)

```
In [11]: sum(networkx.triangles(ring_network(100)).values())/3.0
```

```
Out[11]: 100.0
```

```
In [12]: sum(networkx.triangles(add_shortcuts(ring_network(100))).values())/3.0
```

```
Out[12]: 105.0
```

```
In [13]: def additional_triangles2(N):  
          return sum(networkx.triangles(add_shortcuts(ring_network(N))).values())/3.0 -  
          sum(networkx.triangles(ring_network(N)).values())/3.0
```

```
In [14]: additional_triangles2(100)
```

```
Out[14]: 5.0
```

```
In [15]: additional_triangles2(100000)
```

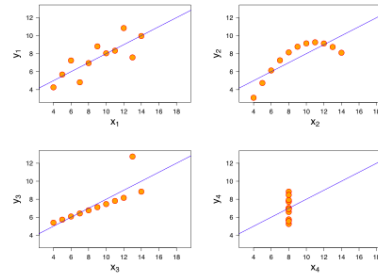
```
Out[15]: 6.0
```

# What causes the discrepancy?


- We are calculating an average, which is doing something strange. A important lesson about statistical quantities:

# Anscombe's quartet

- 11 (x,y) datapoints
- Mean & variance of x match exactly
- Mean & variance of y, correlation between x & y, fitted line (linear regression) match to at least 2 decimal places



# Even better: Datasaurus




<https://www.autodeskresearch.com/publications/samestats>  
(<https://www.autodeskresearch.com/publications/samestats>)




# Let's look at what we are averaging!


```
In [16]: _ = plt.hist(list(networkx.clustering(add_shortcuts(ring_network(100))).values()),  
_ = plt.hist(list(networkx.clustering(add_shortcuts(ring_network(100))).values()),  
bins=100)
```



```
In [17]: _ = plt.hist(list(dict(networkx.degree(add_shortcuts(ring_network(100))))).values  
()), bins=100)
```



```
In [18]: _ = plt.hist(list(dict(networkx.degree(add_shortcuts(ring_network(10000))))).values  
()), bins=100)
```



# Are we sure this explains the discrepancy?

- There is some difference between the two methods, but there might be others
- How can we check? → Homework assignment

# Calculating at a fixed $N$

- Given  $N$ , what is the average clustering coefficient of these networks?

One possible way to calculate: generate a single network, and use that:

```
In [19]: networkx.average_clustering(add_shortcuts(ring_network(10000)))
```

```
Out[19]: 0.3421695238095362
```

Obviously, we shouldn't use a single network for the estimate, but re-run it a few times:

```
In [20]: networkx.average_clustering(add_shortcuts(ring_network(10000)))
```

```
Out[20]: 0.3423485064935197
```

```
In [21]: networkx.average_clustering(add_shortcuts(ring_network(10000)))
```

```
Out[21]: 0.341968686868699
```

Even better, we can calculate an average of these:

```
In [22]: values = [networkx.average_clustering(add_shortcuts(ring_network(10000))) for i in range(10)]
```

```
In [23]: def avg(values):  
         return sum(values) / len(values)
```

```
In [24]: avg([networkx.average_clustering(add_shortcuts(ring_network(10000))) for i in range(10)])
```

```
Out[24]: 0.34188237373738656
```

And re-run this average calculation, as well:

```
In [25]: avg([networkx.average_clustering(add_shortcuts(ring_network(10000))) for i in range(10)])
```

```
Out[25]: 0.34203490620491894
```

```
In [26]: avg([networkx.average_clustering(add_shortcuts(ring_network(10000))) for i in range(10)])
```


```
Out[26]: 0.3416269784104916
```

Of course, then we have the same problem as above -- do we average these averages? And then average those as well?

A better approach: show how our estimate changes as we use larger and larger number of networks -- this is essentially the same as we did for estimating for  $N \rightarrow \infty$ .

```
In [27]: S_values = [3,5,10,15,20,50]
values = [avg([networkx.average_clustering(add_shortcuts(ring_network(10000))) for
i in range(S)]) for S in S_values]
plt.plot(S_values, values)
```

```
Out[27]: [<matplotlib.lines.Line2D at 0x7fd90b419278>]
```




Note that there is an even better way to do this: we don't need to generate a lot of new networks for every datapoint. Instead, we generate a single list of values, and then take successive averages of the first  $n$  elements:

```
In [28]: values = [networkx.average_clustering(add_shortcuts(ring_network(10000))) for i in range(100)]
```

```
In [29]: averages = [avg(values[:S]) for S in range(3, len(values))]
```

```
In [30]: plt.plot(averages)
```

```
Out[30]: [<matplotlib.lines.Line2D at 0x7fd91181c828>]
```



# An important aspect about the sample sizes we might need:


- plotting the histogram of the values for different  $N$  shows some interesting details:

```
In [31]: values_100 = [networkx.average_clustering(add_shortcuts(ring_network(100))) for i in range(1000)]
```

```
In [32]: values_50 = [networkx.average_clustering(add_shortcuts(ring_network(50))) for i in range(1000)]
values_25 = [networkx.average_clustering(add_shortcuts(ring_network(25))) for i in range(1000)]
values_10 = [networkx.average_clustering(add_shortcuts(ring_network(10))) for i in range(1000)]
```



```
In [33]: _ = plt.hist([values_100, values_50, values_25, values_10], bins=50, histtype='step')
```



For each  $N$ , the histogram shows a bell-shaped curve, which is wider for smaller  $N$  and much narrower for larger  $N$ . The center of the peak moves to the left as  $N$  increases.

Of these, the change in the width is a very general feature: most network quantities behave this way: using a larger network decreases the variance, which also lowers the required sample size (number of networks) needed for achieving a given accuracy. The movement of the peak is specific to this problem, and this is the same as we have seen when plotting the convergence as a function of  $N$  before.

# Some interesting features of the random module

We have used or mentioned:

```
In [7]: import random  
        random.shuffle()  
        random.choice()
```

```
In [ ]: But there is also:
```

```
In [8]: random.seed()
```

# An example for using random.seed():

```
In [9]: random.seed(42)  
print(random.random())  
print(random.random())  
print(random.random())
```

```
0.6394267984578837  
0.025010755222666936  
0.27502931836911926
```

```
In [10]: random.seed(42)  
print(random.random())  
print(random.random())  
print(random.random())
```

```
0.6394267984578837  
0.025010755222666936  
0.27502931836911926
```

Which means that we can generate the same random numbers again -- does this make sense?

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

From <https://xkcd.com/221/> (<https://xkcd.com/221/>)

**John von Neumann: "Anyone who attempts to generate random numbers by deterministic means is, of course, living in a state of sin."**

# Having random.seed()

- is it a problem?
- what is it good for?
- PRNG: Pseudo Random Number Generator

- PRNG has internal state, updating that internal state every time it is used
- Starting point of internal state is set by seed()
- PRNG shouldn't be used for cryptography -- attacker might guess internal state, breaking the encryption
- Why can we use a PRNG for network analysis / simulation?

# What do we actually need?

- Want to calculate average over a set of objects
- Can't use all objects (like "networks of size N"), need to use a small sample to approximate
- As long as sample is unbiased, it is fine if it is deterministic



## Also: `random.getstate()`, `random.setstate()`

- Can't use `random.seed()` to get the current state of the PRNG -- need `.getstate()` for that
- Note that `random.getstate()` returns a much larger state than a hash of an object, which `seed()` uses to initialize if not a number is passed in
- `seed()` initializes in a small sub-space of the total possible states of the PRNG

```
In [12]: random.getstate()
```

```
Out[12]: (3,  
          (2468570525,  
           44967195,  
           2667364560,  
           2449893699,  
           1652692239,  
           766678126,  
           273175325,  
           1513475390,  
           2407048223,  
           2326550691,  
           3055735416,  
           2487780036,  
           476975371,  
           81632736,  
           1598452444,  
           3338301038,  
           3898475993,  
           1749546629,  
           4084786842,  
           949316744,  
           2086501466,  
           4175211502,  
           3792229788,  
           1718685282,  
           2499662139,  
           4222931543,  
           3063257123,  
           910424605,  
           1400804300,  
           830603822,  
           3216023045,  
           2756927633,  
           3684278863,  
           3724968901,  
           332416530,  
           52016619,
```

































